


# DATABASE INTRO

JOHN JERNIGAN 8/17/2020



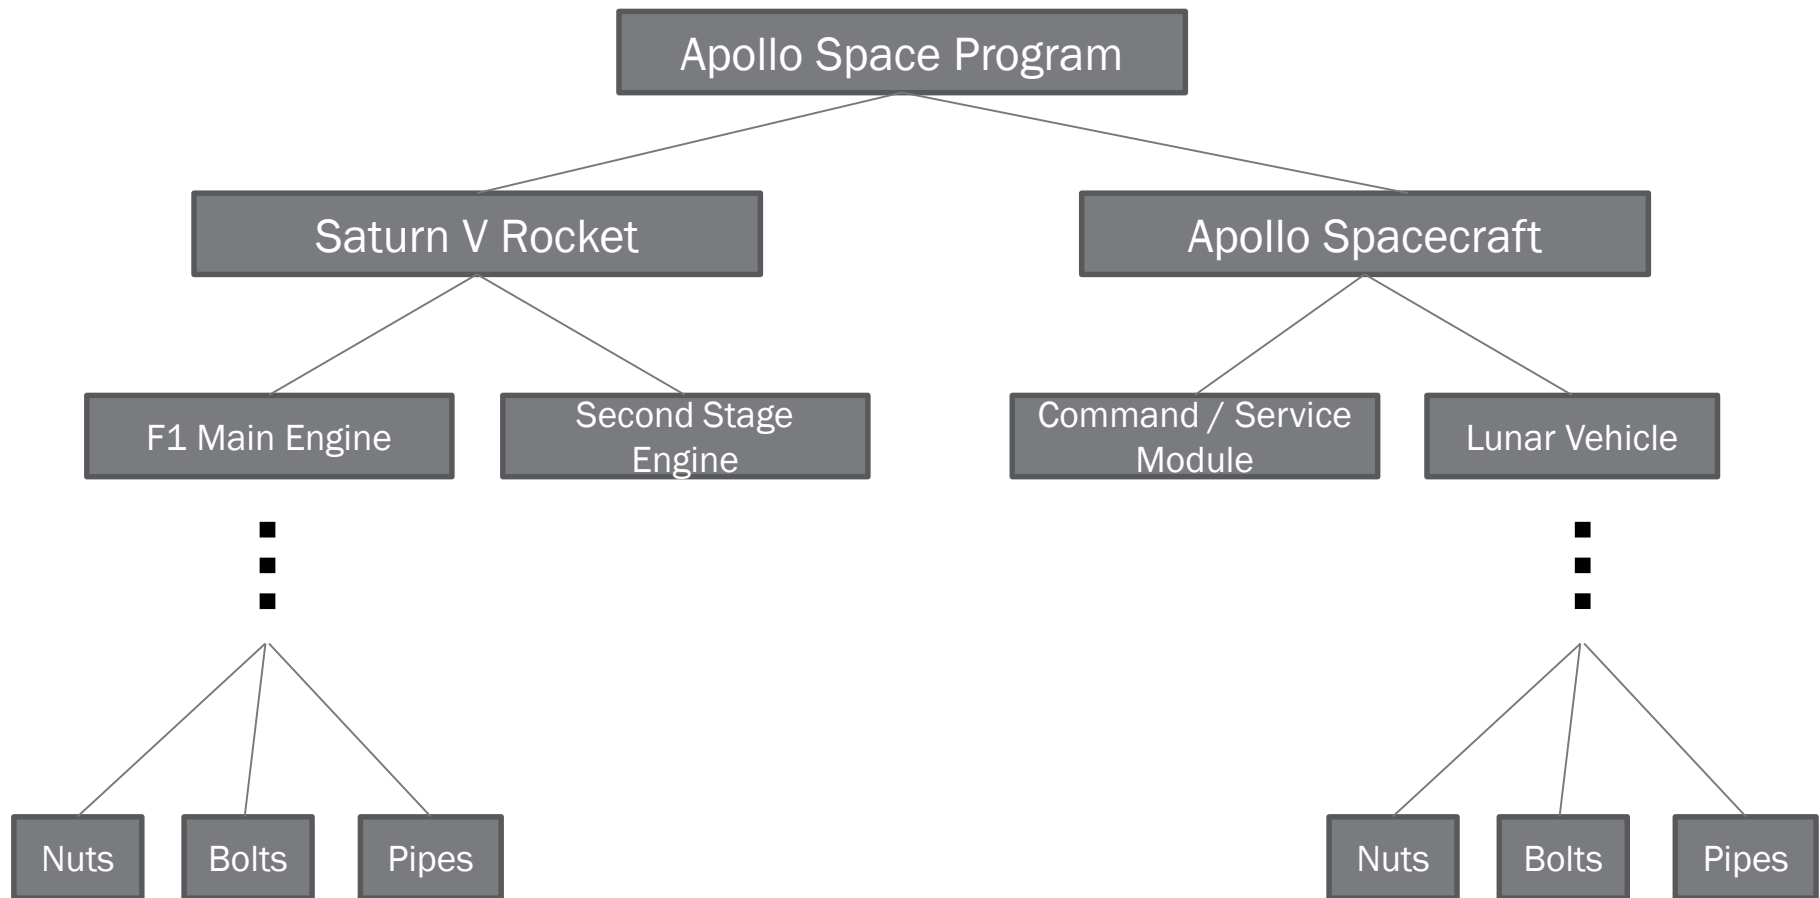


# Information Management System

3 million parts  
inventoried!

1966: First Database Management System (DBMS)

# IMS had a hierarchical “tree” structure, internally



Tree structures are fast and efficient!

# Hierarchical Database Disadvantages:

- Developers **MUST** understand internal database structure
- Which means you really don't want to fundamentally alter the database
- Because that will break your existing applications that use it

# Hierarchical Database Disadvantages:

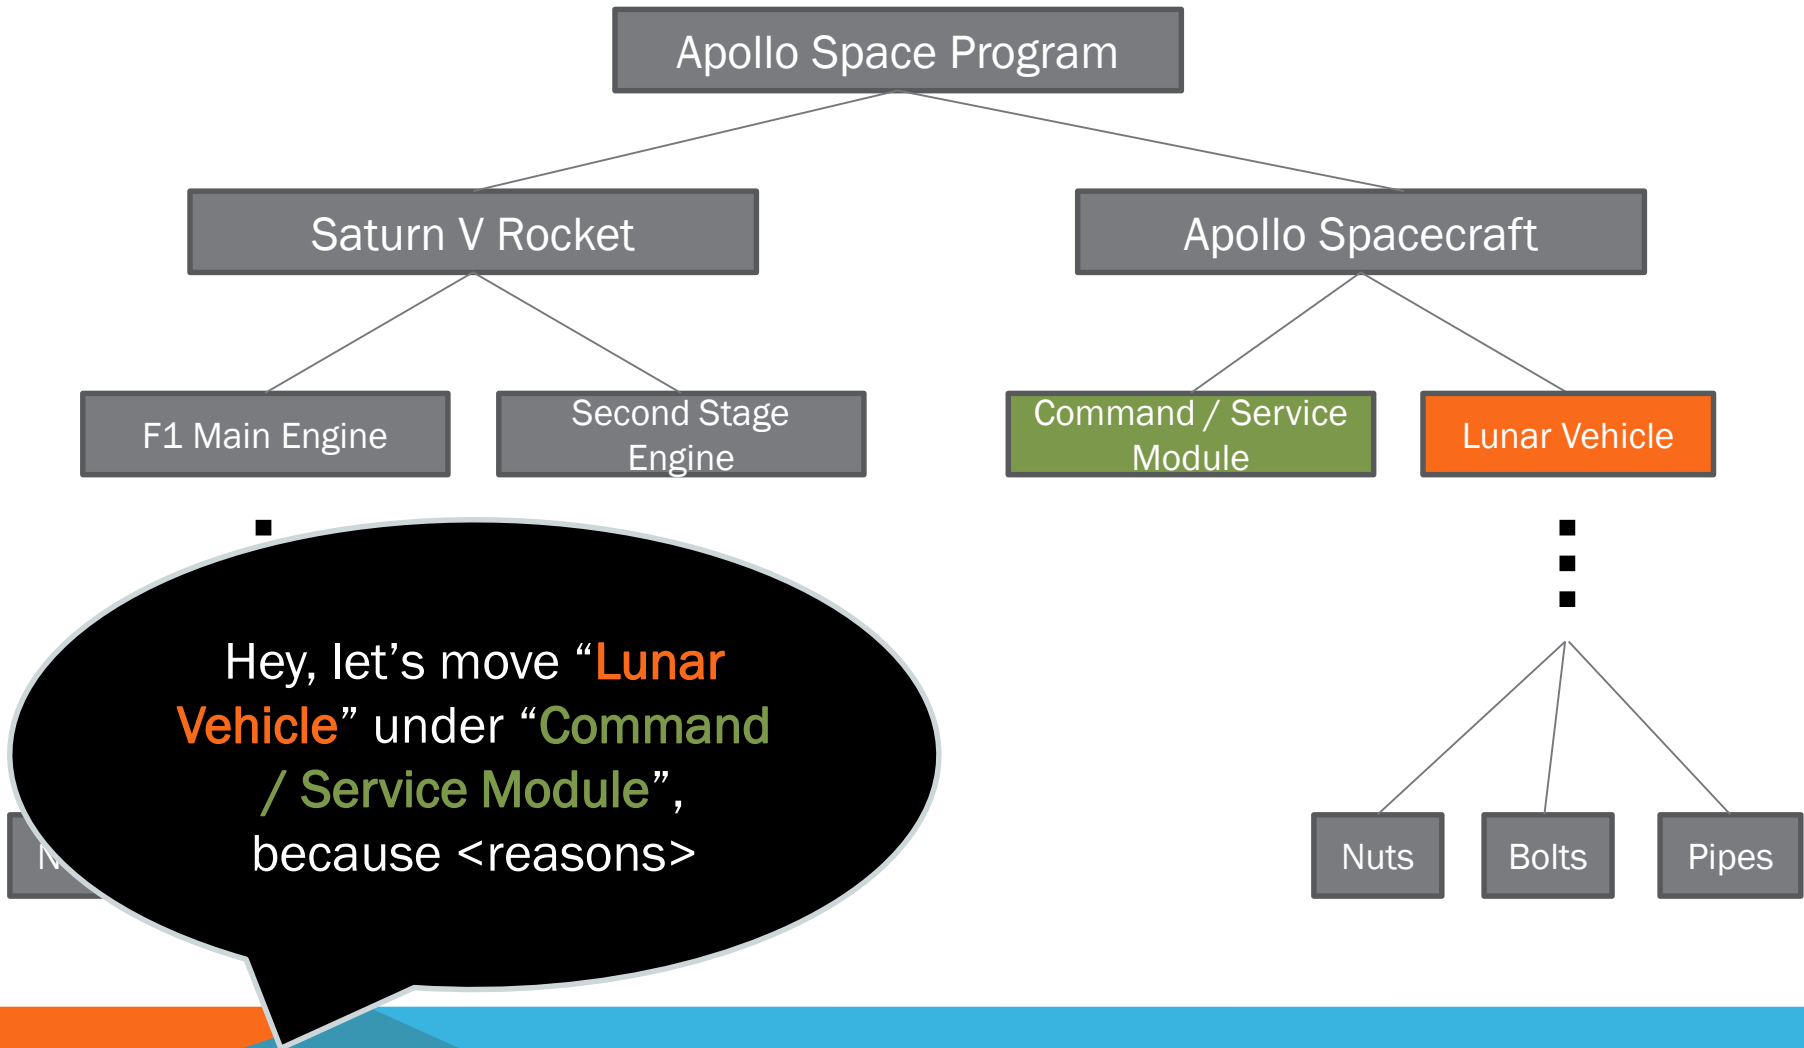
- Developers MUST understand internal database structure
- Which means you really don't want to fundamentally alter the database
- Because that will break your existing applications that use it

# Hierarchical Database Disadvantages:

- Developers MUST understand internal database structure
- Which means you really don't want to fundamentally alter the database
- Because that will break your existing applications that already use it



# You Really Don't Want to Fundamentally Alter a Hierarchical Database






# Hierarchical Database Disadvantages:

- Developers MUST understand internal database structure
- Which means you really don't want to fundamentally alter the database
- Because that will break your existing applications that use the database

# You Really Don't Want to Fundamentally Alter a Hierarchical Database



But ***what if*** there were a database model...

But ***what if*** there were a database model...


...where expensive database  
applications ***continued to work...***

But ***what if*** there were a database model...

...where expensive database  
applications ***continued to work...***

***...even if the internal database  
structure were altered?***

In 1970, Dr. Edgar Codd invents:  
“*Relational* model for database management”



# A Relational Model of Data for Large Shared Data Banks

E. F. Codd

*IBM Research Laboratory, San Jose, California*

Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation). A prompting service which supplies such information is not a satisfactory solution. Activities of users at terminals and most application programs should remain unaffected when the internal representation of data is changed and even when some aspects of the external representation



# Relational Databases Use Table Structures, Not Trees

**Schema**

**Attributes (Columns)**

Unity ID	Name	Email	Phone
jajerni2	John	jajerni2@ncsu.edu	919.513.1666
bwbarbou	Brandon	bwbarbou@ncsu.edu	919.515.0706
<u>avillan</u>	Andrea	<u>avillan@ncsu.edu</u>	919.515.7106

**Tuple (Row)**

# Use Structured Query Language (SQL) to Interface with Database

Unity ID	Name	Email	Phone
jajerni2	John	<a href="mailto:jajerni2@ncsu.edu">jajerni2@ncsu.edu</a>	919.513.1666
bwbarbou	Brandon	<a href="mailto:bwbarbou@ncsu.edu">bwbarbou@ncsu.edu</a>	919.515.0706

```
mysql> SELECT Name,Email FROM staff WHERE 'Unity ID'='jajerni2';
```

```
+-----+-----+
| Name   | Email                               |
+-----+-----+
| John   | jajerni2@ncsu.edu |
+-----+-----+
```

# Creating the First Relational Databases...



# Creating the First Relational Databases...

Ed Oates

Bob Miner

Larry Ellison




# Creating the First Relational Databases...


The Oracle logo is centered within a large red square. The word "ORACLE" is written in a white, bold, sans-serif font. A registered trademark symbol (®) is located at the top right of the letter "E".

ORACLE®

# Some Top Relational Database Management Systems


- **Oracle**
  - Microsoft SQL Server
  - MySQL (free, open-source; see: MariaDB)
  - Postgres (free, open-source)
- 

# Some Top Relational Database Management Systems


- Oracle
  - **Microsoft SQL Server**
  - MySQL (free, open-source; see: MariaDB)
  - Postgres (free, open-source)
- 



# Some Top Relational Database Management Systems

- Oracle
  - Microsoft SQL Server
  - **MySQL (free, open-source; see: MariaDB)**
  - Postgres (free, open-source)
- 

# Some Top Relational Database Management Systems

- Oracle
  - Microsoft SQL Server
  - MySQL (free, open-source; see: MariaDB)
  - **PostgreSQL (free, open-source)**
- 

## Special Mention of Popular Database:



- Stripped-down
- Bare-essentials
- Still powerful (also: free)

# What's Wrong with Relational Databases?



# What's Wrong with Relational Databases?



- Not much

# What's Wrong with Relational Databases?



- Not much
- Until you start dealing with Big Data

# CLASSIC “BIG DATA” DEFINITION



What Can Go Wrong With Relational Databases:  
Big Data



# Database Too Slow?

## Scale Resources **Vertically**



What Can Go Wrong With Relational Databases:  
**Poor Performance**

Database Too Slow?

# Scale Resources Vertically



What Can Go Wrong With Relational Databases:  
**Poor Performance**


Database Too Slow?

**Scale Resources Vertically**



What Can Go Wrong With Relational Databases:  
**Poor Performance**

Or ... **Scale** Resources **Horizontally**




What Can Go Wrong With Relational Databases:  
**Poor Performance**

Or ... **Scale** Resources **Horizontally**



What Can Go Wrong With Relational Databases:  
**Poor Performance**

## Or ... **Scale Resources Horizontally**



What Can Go Wrong With Relational Databases:  
**Poor Performance**

**Scaling Vertically:** Limited by physics, and  
state-of-the-art

Scaling Horizontally: Works! But ...  
presents new problems

What Can Go Wrong With Relational Databases:  
**Poor Performance**

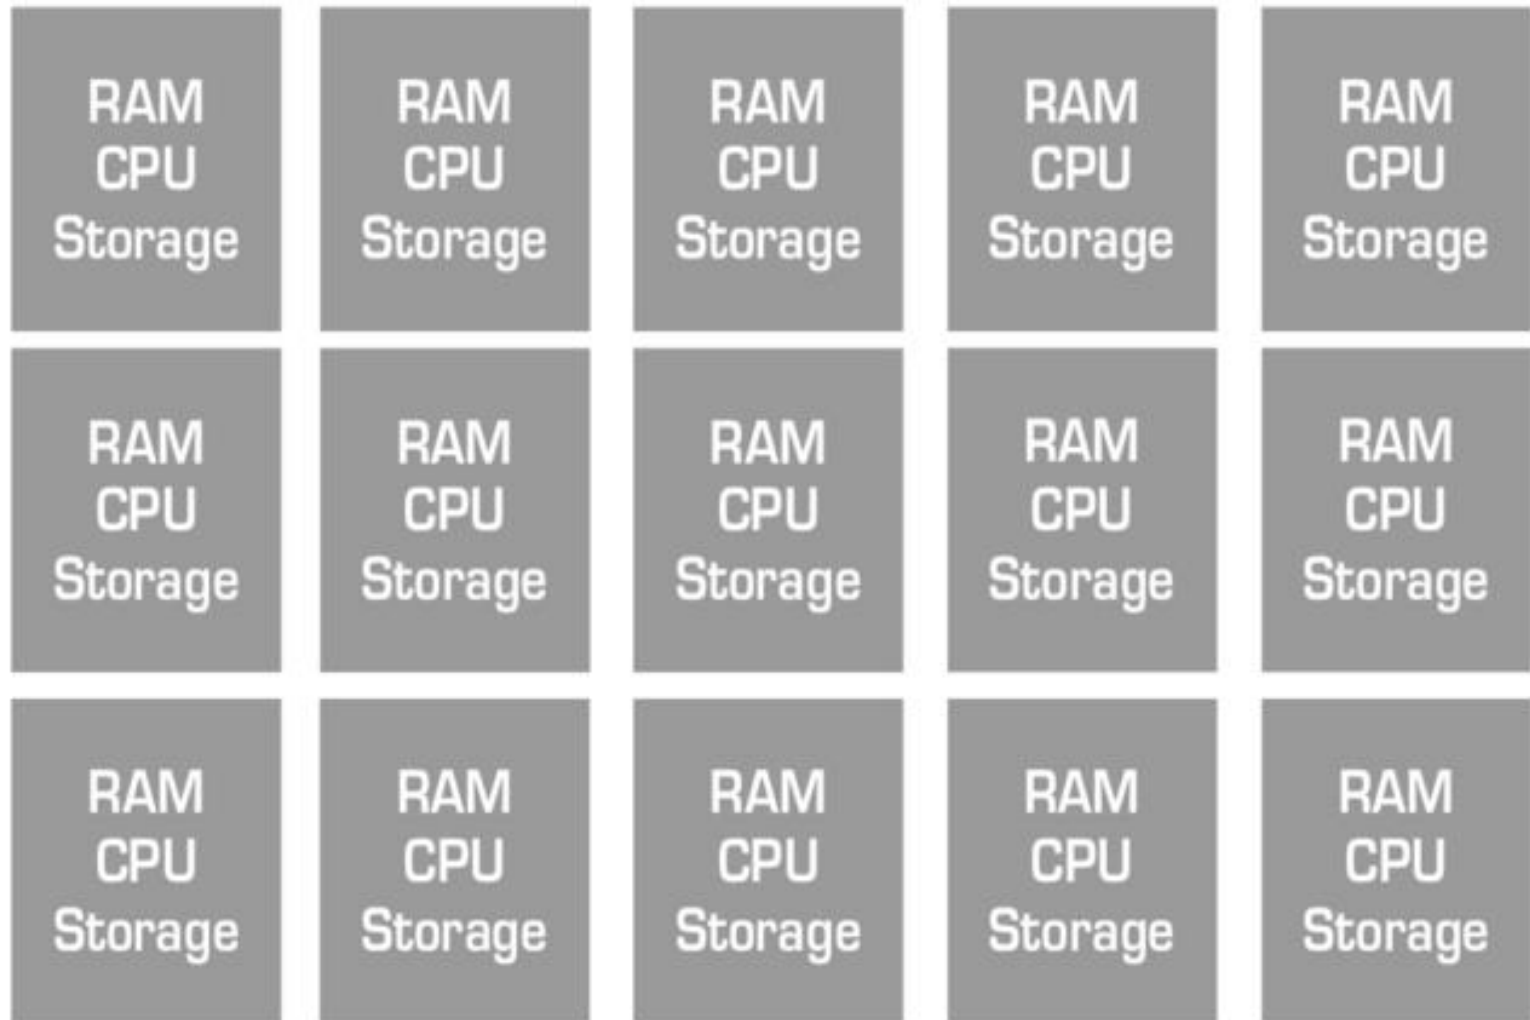


**Scaling Vertically:** Limited by physics, and  
state-of-the-art

**Scaling Horizontally:** Works! But ...  
presents new problems

What Can Go Wrong With Relational Databases:  
**Poor Performance**

Consider the synchronization issues when updating the database simultaneously on multiple nodes



What Can Go Wrong With Relational Databases:  
**Poor Performance**

Engineers attempted to build  
better performing databases.  
They were called...



# NoSQL - A misleading term

Have you heard the term NoSQL?



# NoSQL - A misleading term

NoSQL?

It *seems* to mean  
“not a relational database”

NoSQL - A misleading term

A better way of interpreting it:



Not Only SQL

# NoSQL - A misleading term

Think: Solving Big Data database challenges  
with *application-specific solutions*.



# NoSQL - A misleading term

Think: Solving Big Data database challenges  
with *application-specific solutions*.

Working with big JSON data?  
Try:





# NoSQL - A misleading term

Think: Solving Big Data database challenges  
with *application-specific solutions*.

Working with big key-value pairs?  
Try:



# NoSQL - A misleading term


Think: Solving Big Data database challenges  
with *application-specific solutions*.

Working with big graph data?  
Try:



# NoSQL - A misleading term

Think: Solving Big Data database challenges  
with *application-specific solutions*.




# Recap: The Big Picture of Relational Databases and NoSQL




# Recap: The Big Picture of Relational Databases and NoSQL



**IBM IMS**



**Relational** Database  
Management Systems  
(RDBMS)



**NoSQL** Database  
Management Systems





# Recap: The Big Picture of Relational Databases and NoSQL

## Gartner “Strategic Planning Assumption”:


By 2017, all leading operational DBMSs will offer multiple data models, relational and NoSQL, in a single DBMS platform.



RDBMS



NoSQL



RDBMS + NoSQL =  
*Flying Car*



# **PRACTICUM TIP: USE INDEXING**

**If you have a SAS7BDAT dataset...**





# **PRACTICUM TIP: USE INDEXING**

If you have a SAS7BDAT dataset...

**Or your data is inside MySQL / Postgres / MSSQL...**



# PRACTICUM TIP: USE INDEXING

If you have a SAS7BDAT dataset...

Or your data is inside MySQL / Postgres / MSSQL...

**You NEED to CONSIDER where to  
create indexes for performance purposes.**

## PRACTICUM TIP: USE INDEXING

✓ *Dramatically* decrease query times!

E.g. 1 hour → 2.4 seconds

I have seen this... that is a 1500x speedup.

## PRACTICUM TIP: USE INDEXING

- ✓ *Dramatically* decrease query times!
- ✓ *Speed up* time to sort data!

E.g. 1 hour → 2.4 seconds

I have seen this... that is a 1500x speedup.

# PRACTICUM TIP: USE INDEXING

- Indexes are created per-column (*attribute*)

e.g.

- Employee Number
- Purchase Price
- Transaction Date

Index them all!

## PRACTICUM TIP: USE INDEXING

- Indexes are created per-column (*attribute*)
- Indexes can hugely speed up your queries!

# PRACTICUM TIP: USE INDEXING

- Indexes are created per-column (*attribute*)
- Indexes can hugely speed up your queries!
- Index updates are performance overhead when data is added to database which is a tradeoff to consider.

MSA Students:

Don't worry about this!

Real-World Production Systems: Worry about this!

## Consider Which Columns You are Querying Against

E.g. will you be searching the table by Name?  
Unity ID? All columns?

**Schema**

**Attributes (Columns)**

Unity ID	Name	Email	Phone
jajerni2	John	<a href="mailto:jajerni2@ncsu.edu">jajerni2@ncsu.edu</a>	919.513.1666
bwbarbou	Brandon	<a href="mailto:bwbarbou@ncsu.edu">bwbarbou@ncsu.edu</a>	919.515.0706
<u>avillan</u>	Andrea	<a href="mailto:avillan@ncsu.edu">avillan@ncsu.edu</a>	919.515.7106

**Tuple (Row)**



## Demo: Understand Database Indexing



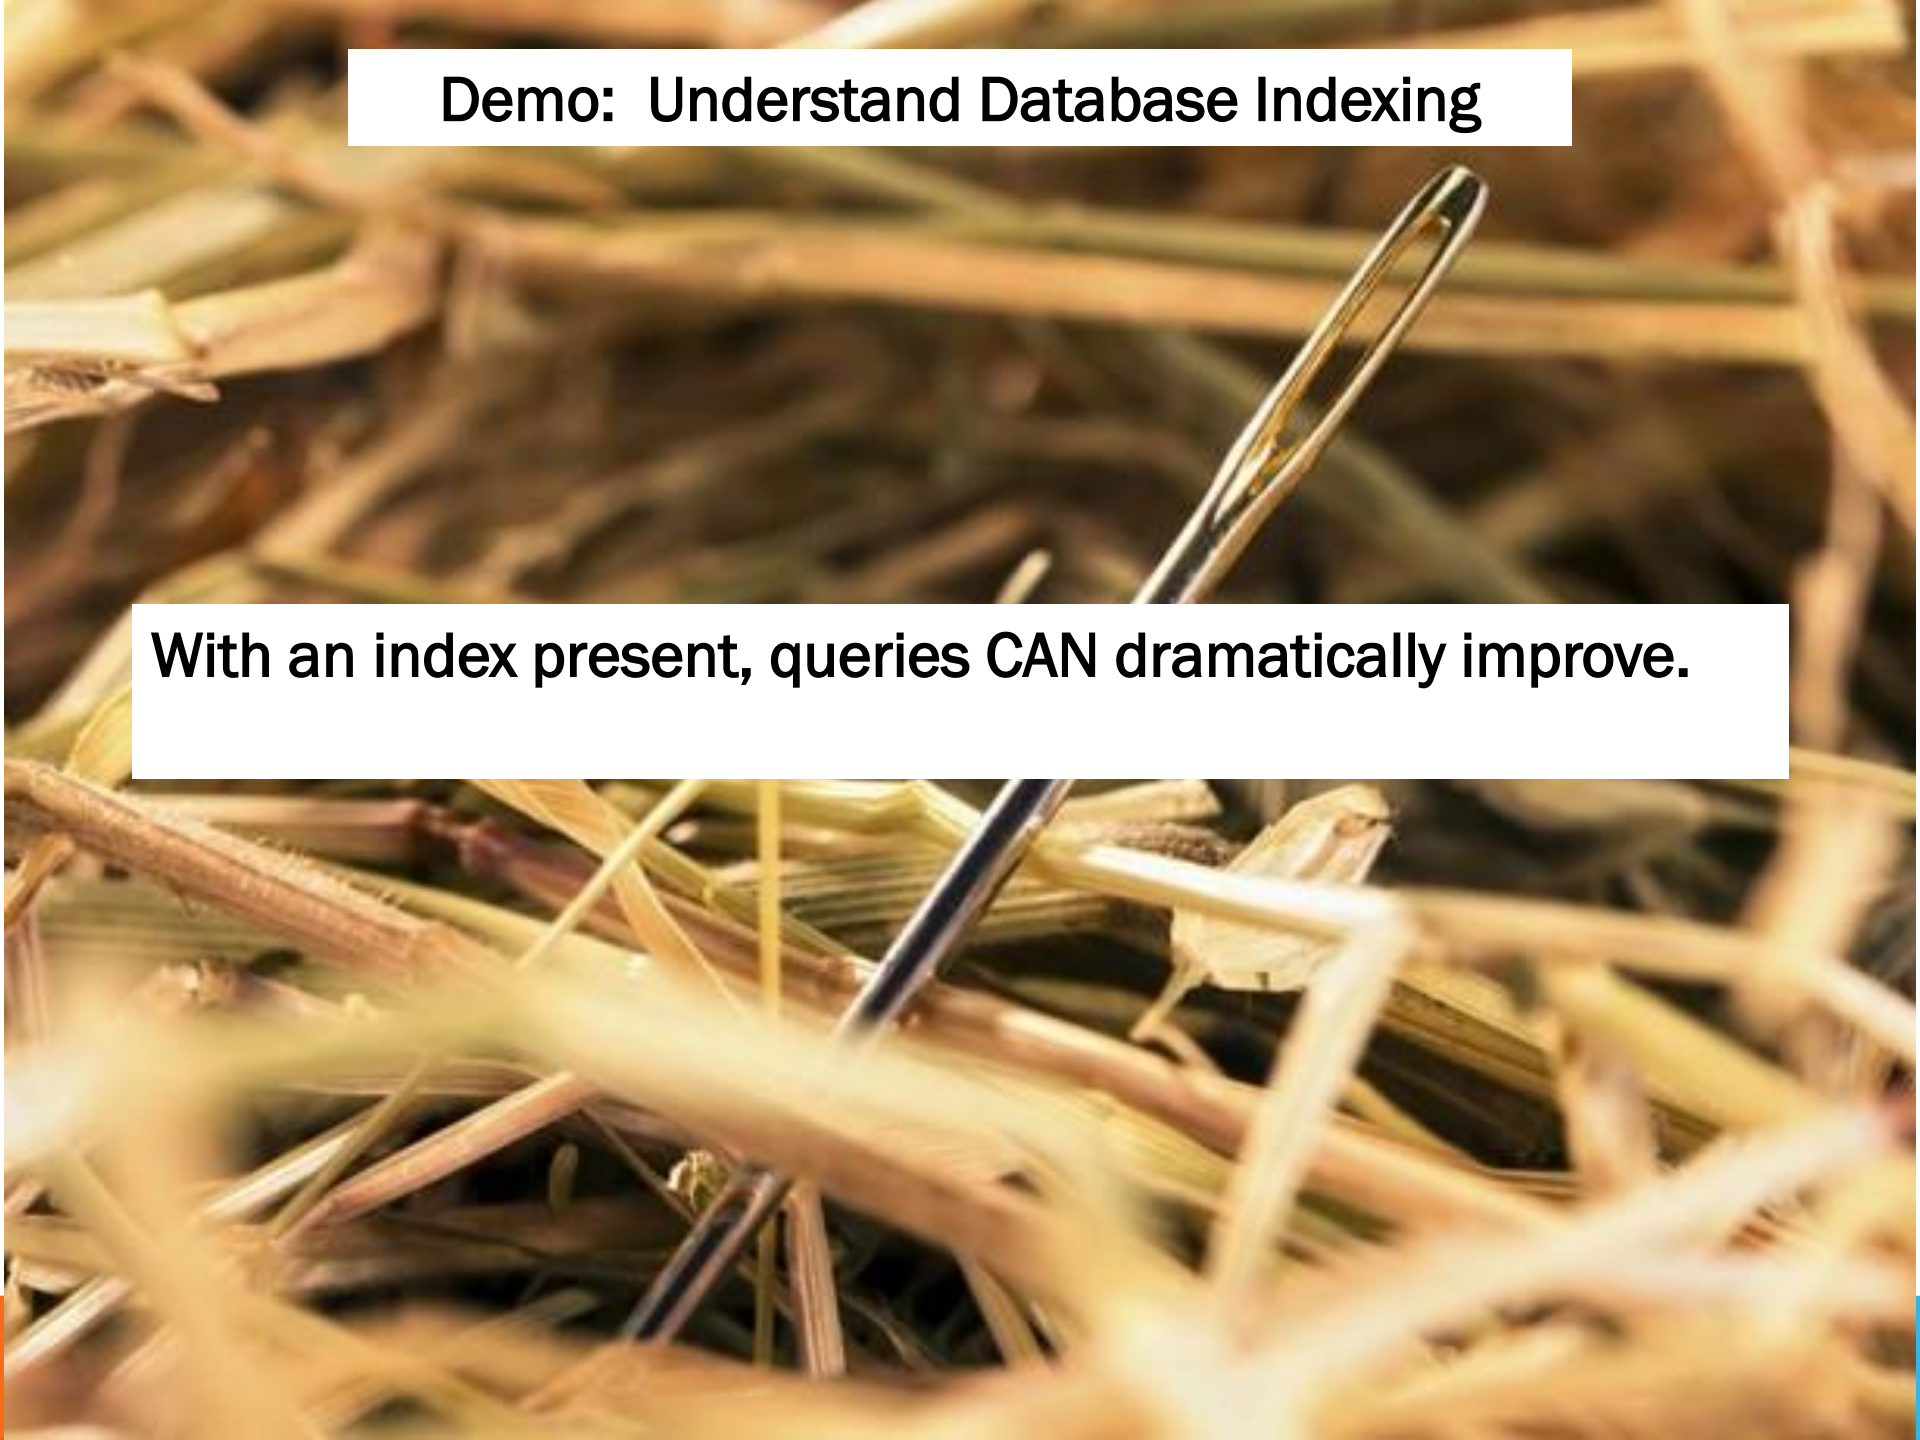
## Demo: Understand Database Indexing





## Demo: Understand Database Indexing

With an index present, queries CAN dramatically improve.



## Demo: Understand Database Indexing

Indexes help with “finding a needle in a haystack”.

They do not help (and do not hurt) if you're going to extract most of the database anyway:

```
SELECT COUNT(Unity_ID) FROM directory;
```

(this returns the number of NCSU students in a directory)

## Index Your Database Columns!

**As a data scientist...**

**... if you use SQL queries in a database...**

**As a data scientist...**

**... if you use SQL queries in a database...**

**... which includes proc sql on a SAS7BDAT file ...**

**... or SQL with PostgreSQL ...**

## Index Your Database Columns!

As a data scientist...

... if you use SQL queries in a database...

... which includes proc sql on a SAS7BDAT file ...

... or SQL with PostgreSQL ...

... learn how to turn on indexing for any relevant data columns.

## Index Your Database Columns!

As a data scientist...

- ... if you use SQL queries in a database...

  - ... which includes proc sql on a SAS7BDAT file ...

  - ... or SQL with PostgreSQL ...

... learn how to turn on indexing for any relevant data columns.

(e.g. index “date” columns for transactional data)



## Index Your Database Columns!

As a data scientist...

... if you use SQL queries in a database...

... which includes proc sql on a SAS7BDAT file ...

... or SQL with PostgreSQL ...

... learn how to turn on indexing for any relevant data columns.

(e.g. index “date” columns for transactional data)

**This will save you time** whenever you’re looking for a small bit of data in a large dataset...

... but it’s up to YOU to turn on indexing.

# For more info on indexing...

- **SAS Indexes:** <http://www2.sas.com/proceedings/sugi29/123-29.pdf>  
Also: SAS Programming 3 book, section 3-5
- **SQLite Indexes:** <http://www.sqlitetutorial.net/sqlite-index/>
- **PostgreSQL Indexes:**  
<https://www.postgresql.org/docs/current/static/indexes.html>
- **General Info:** [https://en.wikipedia.org/wiki/Database\\_index](https://en.wikipedia.org/wiki/Database_index)

# Questions

